

Paul Wendt

Email: paulwendt567@gmail.com | Phone: +1-609-635-6144
LinkedIn: [linkedin.com/in/paul-wendt-81380a260](https://www.linkedin.com/in/paul-wendt-81380a260) | GitHub: github.com/WillChangeThisLater

Summary

Senior Infrastructure and Data Engineer with five years of experience designing, migrating, and optimizing scalable cloud infrastructure. Experienced in building pipeline orchestration frameworks, managing distributed compute workloads for machine learning, and writing reliable data platform primitives.

Experience

SimpliSafe — Boston, MA

Senior ML Ops Engineer

Oct 2024 - Mar 2026

Owned the data infrastructure feeding computer vision model training.

- Built and operated four production data services (video ingestion, AI-permission routing, Right-to-be-Forgotten deletion) on EventBridge → SQS → Lambda → S3/Kinesis Firehose
- Built Ray/Anyscale pipelines moving nightly video batches and metadata JSON between S3 and an annotation platform, orchestrated via MWAA Airflow
- Deployed Voxel51 on AWS EKS, establishing secure connectivity to MongoDB Atlas and engineering a custom Anyscale operator adopted by the vendor as a core product feature
- Ran a 4-month enablement program teaching 3 engineers CloudFormation, CI/CD, and on-call triage until they owned the video-ingestion service independently

Data Engineer II

Nov 2021 - Oct 2024

Major implementer on the team that migrated the data platform from cron jobs on EC2 to Dagster orchestration.

- Migrated legacy EC2 cron pipelines to an orchestrated source/transform/sink model on Dagster, scaling the platform to thousands of daily runs
- Co-authored a YAML-based DSL on top of the new architecture to democratize pipeline creation for 30+ analysts across 10 teams
- Designed a custom Dagster dependency sensor to preserve partition lineage across 350+ downstream pipelines
- Migrated GitHub Actions runners from a standalone EC2 box to EKS (CDK + Helm), cutting integration test runtimes from over an hour to under 20 minutes

John Hancock Life Insurance — Boston, MA

Actuarial Associate

May 2019 - Nov 2021

- Implemented annuity valuation model with reserving team and state regulators; wrote a Python AST transpiler converting model specifications to Excel, eliminating weeks of manual translation effort
- Automated ~5 core quarterly reporting workbooks with VBA, cutting ~10 hours per quarter from the reporting cycle

Education

Temple University — Philadelphia, PA

Aug 2015 - May 2019

B.A. in Actuarial Science, Minor in Computer Science · *Summa Cum Laude*

Certifications

- Certified Kubernetes Administrator (CKA)** — 2023
- AWS Certified Cloud Practitioner** — 2022
- Associate of the Society of Actuaries (ASA)** — 2021

Projects

- Trail PinePhone Build:** Configured a mobile Linux environment utilizing Tailscale VPN and SSH-over-hotspot to establish secure, remote shell sessions into a distributed Anyscale cluster during a 4.5 month thru-hike
- AI Job Search Pipeline:** Multi-agent pipeline that sweeps job boards daily, vets postings against a structured evidence bank, and suggests the best listings
- Pi Harness Extension:** Extended open-source Pi harness to support multi-modal audio/video input streaming, push-to-talk dictation (STT), and custom TUI model-state indicators

Technical Skills

Category	Technologies
Languages	Python, SQL, Bash, Go
Data	Apache Iceberg, Dagster, Kafka, Airflow/MWAA, Athena, Pandas, DuckDB, Hex
Cloud & Infra	AWS (EKS, CDK, Lambda, Step Functions, DynamoDB, S3, SQS, Kinesis Firehose, EventBridge), Kubernetes, Helm, Ray/Anyscale
DevOps	GitHub Actions (custom actions, EKS runners), CI/CD architecture, ML data pipelines
Tools	Git, Neovim, Tmux, SSH, jq/yq, GraphQL