

Paul Wendt

Email: paulwendt567@gmail.com | Phone: +1-609-635-6144
LinkedIn: [linkedin.com/in/paul-wendt-81380a260](https://www.linkedin.com/in/paul-wendt-81380a260) | GitHub: github.com/WillChangeThisLater

Summary

Senior ML Ops and Data Engineer with five years of experience designing, migrating, and optimizing scalable cloud infrastructure. Experienced in building pipeline orchestration frameworks, managing distributed compute workloads for machine learning, and writing reliable data platform primitives.

Experience

SimpliSafe — Boston, MA

Senior ML Ops Engineer

Oct 2024 - Mar 2026

Owned the data infrastructure feeding computer vision model training.

- Built and operated four production data services (video ingestion, AI-permission routing, Right-to-be-Forgotten deletion) on EventBridge → SQS → Lambda → S3/Kinesis Firehose
- Designed a version-controlled video sampling framework which cut video ingestion volume ~75% — 1.2M to ~300K videos/day
- Productized data governance as daily Dagster jobs: SQL invariants and S3 inventory reconciliations verified deletion events, detected ingest drift against upstream Kafka state, and caught multiple cross-service bugs
- Built Ray/Anyscale pipelines moving nightly video batches and metadata JSON between S3 and an annotation platform, orchestrated via MWAA Airflow

Data Engineer II

Nov 2021 - Oct 2024

Major implementer on the team that migrated the data platform from cron jobs on EC2 to Dagster orchestration.

- Migrated legacy cron pipelines to a source/transform/sink model on Dagster, contributing to a YAML-based DSL used by ~30 analysts and data engineers across 10 teams; platform scaled to hundreds of pipelines with thousands of daily runs
- Designed a dependency sensor preserving partition lineage where Dagster's default triggered downstream pipelines with the wrong partition keys, which became a core platform primitive across ~350 pipelines
- Migrated GitHub Actions runners from a single EC2 box to EKS (CDK + Helm, least-privilege IAM); parallelized test suites, cutting integration test runtimes from over an hour to under 20 minutes
- Authored ~50 pipelines syncing ~25 Zuora objects into Apache Iceberg tables to power subscription and finance reporting, translating non-technical analysts' requirements into production pipelines

John Hancock Life Insurance — Boston, MA

Actuarial Associate

May 2019 - Nov 2021

- Implemented annuity valuation model with reserving team and state regulators; wrote a Python AST transpiler converting model specifications to Excel, eliminating weeks of manual translation effort
- Automated ~5 core quarterly reporting workbooks with VBA, cutting ~10 hours per quarter from the reporting cycle

Education

Temple University — Philadelphia, PA

Aug 2015 - May 2019

B.A. in Actuarial Science, Minor in Computer Science · *Summa Cum Laude*

Certifications

- Certified Kubernetes Administrator (CKA)** — 2023
- AWS Certified Cloud Practitioner** — 2022
- Associate of the Society of Actuaries (ASA)** — 2021

Projects

- Trail PinePhone Build:** Modified PinePhone for SSH-over-iPhone-hotspot remote development during an Appalachian Trail thru-hike
- AI Job Search Pipeline** (*Python + pi agent harness*): Multi-agent pipeline that sweeps job boards daily, vets postings against a structured evidence bank, and suggests the best listings

Technical Skills

Category	Technologies
Languages	Python, SQL, Bash, Go
Data	Apache Iceberg, Dagster, Kafka, Airflow/MWAA, Athena, Pandas, DuckDB, Hex
Cloud & Infra	AWS (S3, Lambda, SQS, Kinesis Firehose, EventBridge, EKS, CDK), Kubernetes, Helm, Ray/Anyscale
DevOps	GitHub Actions (custom actions, EKS runners), CI/CD architecture, ML data pipelines
Tools	Git, Neovim, Tmux, SSH, jq/yq, GraphQL