

Paul Wendt

Email: paulwendt567@gmail.com | Phone: +1-609-635-6144
LinkedIn: [linkedin.com/in/paul-wendt-81380a260](https://www.linkedin.com/in/paul-wendt-81380a260) | GitHub: github.com/WillChangeThisLater

Summary

Senior Infrastructure and Data Engineer with five years of experience building production systems in regulated environments, plus prior insurance-domain experience at John Hancock. Strong in backend service reliability, compliance controls, and day-to-day AI-assisted engineering with human-in-the-loop validation

Experience

SimpliSafe — Boston, MA

Senior ML Ops Engineer

Oct 2024 - Mar 2026

Owned the data infrastructure feeding computer vision model training.

- Built and operated four production data services (video ingestion, AI-permission routing, Right-to-be-Forgotten deletion) on EventBridge → SQS → Lambda → S3/Kinesis Firehose
- Built Ray/Anyscale pipelines moving nightly video batches and metadata JSON between S3 and an annotation platform, orchestrated via MWAA Airflow
- Deployed Voxel51 on AWS EKS, establishing secure connectivity to MongoDB Atlas and engineering a custom Anyscale operator adopted by the vendor as a core product feature
- Ran a 4-month enablement program teaching 3 engineers CloudFormation, CI/CD, and on-call triage until they owned the video-ingestion service independently

Data Engineer II

Nov 2021 - Oct 2024

Major implementer on the team that migrated the data platform from cron jobs on EC2 to Dagster orchestration.

- Migrated legacy EC2 cron pipelines to an orchestrated source/transform/sink model on Dagster, scaling the platform to thousands of daily runs
- Co-authored a YAML-based DSL on top of the new architecture to democratize pipeline creation for 30+ analysts across 10 teams
- Authored ~50 enterprise data pipelines syncing complex Zuora billing objects via ZOQL and AQuA APIs into Apache Iceberg tables to power core financial and subscription revenue reporting datasets
- Migrated GitHub Actions runners from a standalone EC2 box to EKS (CDK + Helm), cutting integration test runtimes from over an hour to under 20 minutes

John Hancock Life Insurance — Boston, MA

Actuarial Associate

May 2019 - Nov 2021

- Implemented annuity valuation model with reserving team and state regulators; wrote a Python AST transpiler converting model specifications to Excel, eliminating weeks of manual translation effort
- Automated ~5 core quarterly reporting workbooks with VBA, cutting ~10 hours per quarter from the reporting cycle

Education

Temple University — Philadelphia, PA

Aug 2015 - May 2019

B.A. in Actuarial Science, Minor in Computer Science · *Summa Cum Laude*

Certifications

- Certified Kubernetes Administrator (CKA)** — 2023
- AWS Certified Cloud Practitioner** — 2022
- Associate of the Society of Actuaries (ASA)** — 2021

Projects

- Trail PinePhone Build:** Configured a mobile Linux environment utilizing Tailscale VPN and SSH-over-hotspot to establish secure, remote shell sessions into a distributed Anyscale cluster during a 4.5 month thru-hike
- AI Job Search Pipeline:** Multi-agent pipeline that sweeps job boards daily, vets postings against a structured evidence bank, and suggests the best listings
- Pi Harness Extension:** Extended open-source Pi harness to support multi-modal audio/video input streaming, push-to-talk dictation (STT), and custom TUI model-state indicators

Technical Skills

Category	Technologies
Languages	Python, SQL, Bash, Go
Data	Apache Iceberg, Dagster, Kafka, Airflow/MWAA, Athena, Pandas, DuckDB, Hex
Cloud & Infra	AWS (EKS, CDK, Lambda, Step Functions, DynamoDB, S3, SQS, Kinesis Firehose, EventBridge), Kubernetes, Helm, Ray/Anyscale
DevOps	GitHub Actions (custom actions, EKS runners), CI/CD architecture, Git, jq/yq, GraphQL, Tmux, SSH